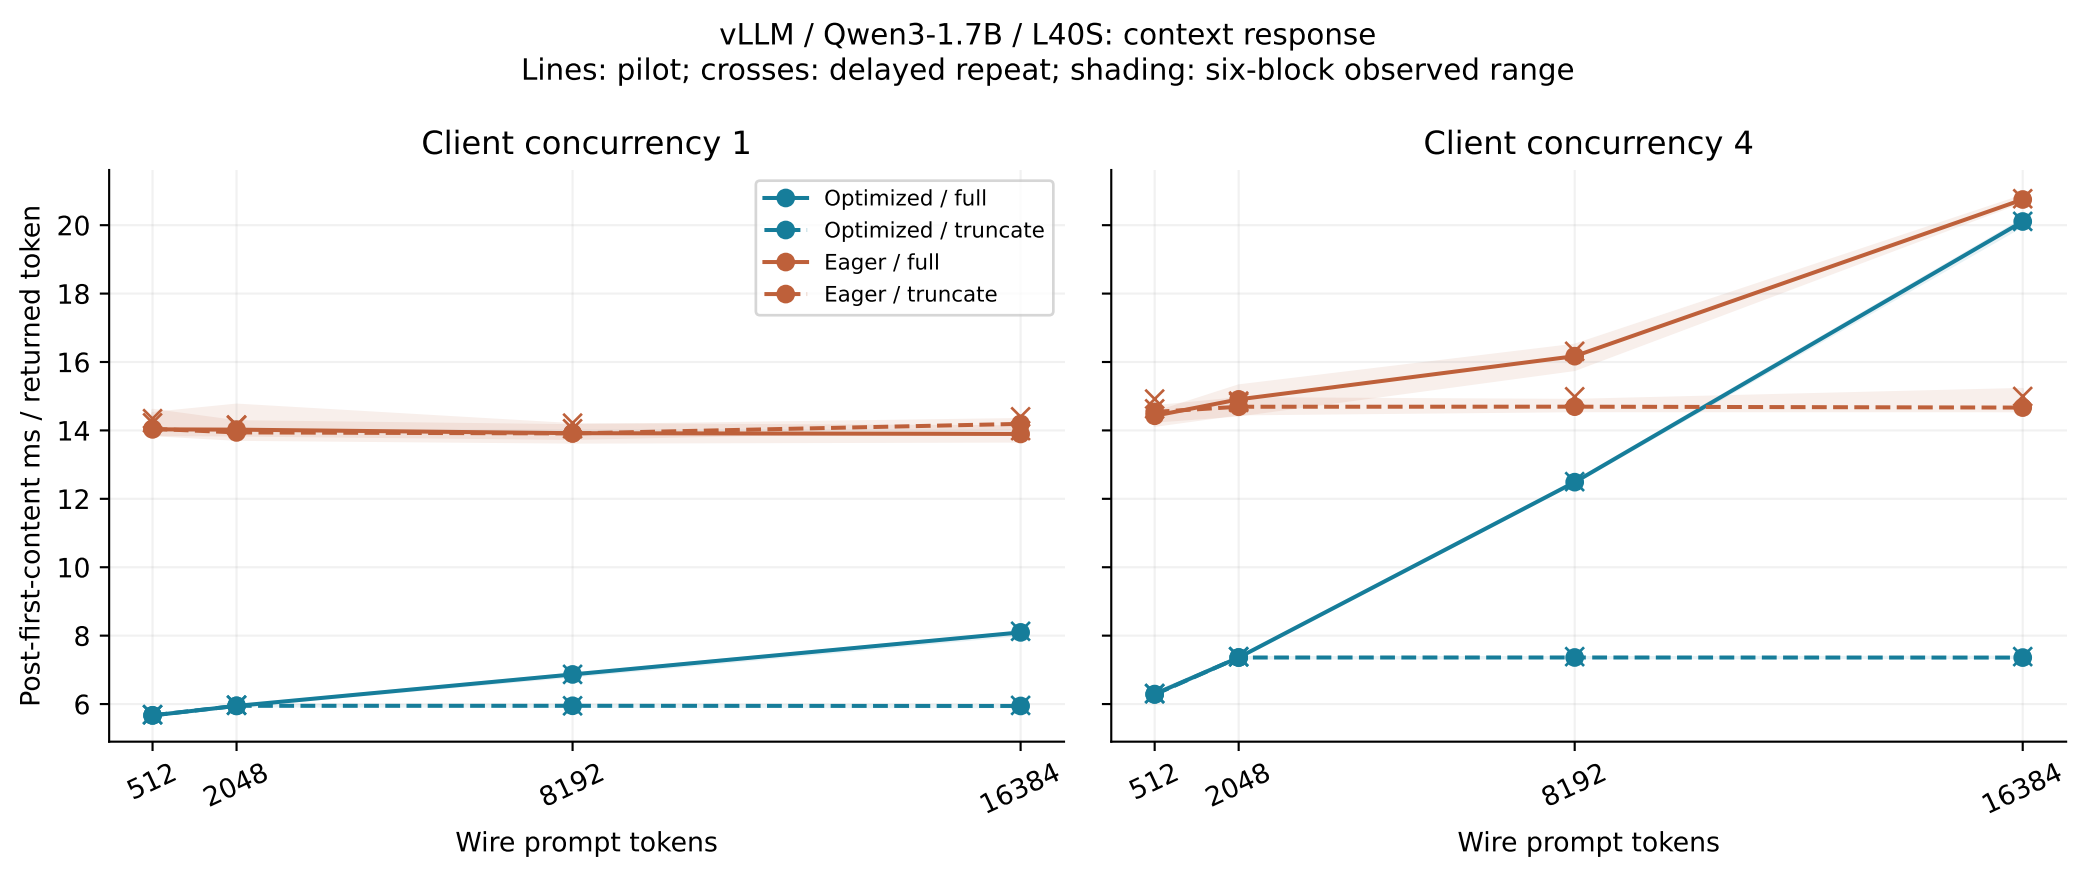
vLLM / Qwen3-1.7B / L40S: context response
Lines: pilot; crosses: delayed repeat; shading: six-block observed range
Client concurrency 1
Client concurrency 4
Optimized / full
Optimized / truncate
Eager / full
Eager / truncate
Post-first-content ms / returned token
20
18
16
14
12
10
8
6
512
2048
8192
16384
Wire prompt tokens