

How do you audit a computation you cannot see?

A treaty clause, a promise to a regulator and an API bill are the same claim. An outsider has two ways to test it: the bytes that came back, and the clock.

Jiawei Li¹ · mentored by Gabriel Kulp² and Roy Rinberg²

¹MATS 10.0 Scholar ²MATS 10.0 Mentor Code + every run artifact: github.com/lijiawei20161002/inference-verification

THE AUDIT PROBLEM

Two governments agree to stop training something. Who checks — and with what?

Every compute-governance instrument ends in a promise about work done inside a building the auditor cannot enter. Nothing in the stack between a GPU and the outside world emits a statement about what the silicon actually did, so the auditor is left with whatever crosses the boundary.




Figure 1 The auditable surface of an inference provider. The same shape recurs one level up — a no-training clause is a promise about which computation ran on a fleet — so the honest place to start is the case where the auditor holds a spec and receives a stream, and every claim can be measured on real hardware.

WHY THE PROMISE IS WORTH BREAKING

A 4-bit copy, a smaller sibling or a truncated context is pure margin, and invisible in the tokens the client pays for. But the deviations worth doing are exactly the ones that **read fewer bytes** — which is why they leave a second trace.

what the provider does instead	bytes it saves	worth it?
Serve Qwen3-0.6B as Qwen3-1.7B	2.83×	capability
4-bit (NF4) weights instead of bf16	2.55×	capacity
Attend half the contracted context	1.35×	capacity
Wrong seed / wrong temperature	1.00×	no

Table 1 Weight + KV bytes per decode step, from the timing grid's own byte counts (clock_channel.json, H100 PCIe, B=1). The last row is the control: a deviation that saves no bytes leaves no clock signature at all — and the two channels split exactly there.

THE TWO CHANNELS ARE BLIND IN OPPOSITE DIRECTIONS

CHANNEL 1 — WHAT IT SAID	CHANNEL 2 — WHEN IT SAID IT
Recompute M on the claimed tokens and read the per-token margin — or skip M and score against a cheap trusted proxy q .	Read the arrival times of a stream already bought. Zero audit FLOPs, zero cooperation from the provider, zero added latency.
sees the sampler, the seed, the exact distribution; wholesale substitution at AUC 0.998 with no M	sees which positions were attended, and which parameters were read at prefll
blind an fp8 KV cache — It moves the served distribution less than M's own rerun noise	blind anything reading the same bytes: wrong seed, wrong temperature, a top-k bug

Figure 2 Complementary by mechanism, not by accident: a deviation either changes the distribution the tokens were drawn from, or the bytes the GPU had to move, or both. One protocol scores both channels — standardized pAUC @ FPR $\leq 0.5\%$, held-out honest calibration, batch/pool ratio $\leq 10\%$ enforced in code.

CHANNEL 1 · READ WHAT CAME BACK

One token proves almost nothing. A verdict is bought in bulk — and the price is fixed before you start.

Re-run M on the tokens the provider claims it produced and compare the per-token sampling margin. The honest and the cheating stream overlap almost completely: the audit rests on one number, the per-token separation d' .




Figure 3 The pooling law with this experiment's own measured inputs — Qwen3-1.7B under 4-bit weights, per-token $d' = 0.075$ (cost_of_a_verdict_scores.npz). Against measured pAUC at 5 pool sizes on a 56,160-token pool the law holds to a mean absolute residual of 0.012 (pool_scaling.json) — which is what lets an auditor price a verdict before spending a budget.

WHAT CATCHES WHAT

Every detector is structurally blind to a different class of deviation, and which one wins is set by the **attack**, not the model. An auditor has to hold a portfolio.

	token DiFR	cross entropy	activation DiFR	token top-loc
4-bit weights	0.718	0.505	1.000	0.600
fp8 KV cache	0.511	0.507	1.000	0.517
wrong temperature	0.516	0.823	0.504	0.661
wrong seed	1.000	0.524	0.575	0.507
top- k bug ($k=2$)	0.553	0.510	0.499	0.523
top- k bug ($k=32$)	0.712	0.605	0.499	0.686

Table 2 headline_ratio.json. Re-scored from the **same** per-token scores as a previously published table that resampled from too small a pool, 16 of 24 cells fall, by a median of 0.137: an over-ratio pool does not exaggerate a signal, it invents one.

WHAT ONE VERDICT COSTS

Detectors are usually ranked by accuracy. An auditor buys neither an AUC nor a FLOP count — it buys tokens. Below: the cheapest honest verdict for each deviation, and what the **same** verdict costs if the provider is made to open up.




Figure 4 cost_of_a_verdict.json; Qwen3-1.7B, $\delta^* = 3.77$, 141 μ s of H100 per recomputed token. Opening the provider up is worth 2,000 $\times$ against quantization and **nothing** against a tampered sampler, which redraws from the honest distribution. Across the whole grid 14 of 35 cells have no price at all — and ranking by AUC hides every one of them.

CHANNEL 2 · READ WHEN IT CAME BACK

A token cannot arrive before its bytes have moved. The saving and the evidence are the same quantity.

One token at a time a GPU is memory-bound: it must drag every weight it claims and read every position it charged for before it can emit anything, so $t_{\text{token}} \geq \text{bytes} / \text{bandwidth}$. Nothing below is drawn — the whole column is timed on one H100.




Figure 5 clock_channel.json, H100 PCIe, B=1, 256-token context, 3 $\times$ 96 timed steps per cell. (A) The 4-bit provider reads 2.55 $\times$ fewer bytes and still costs 1.16 $\times$ what this line prices them at: dequantization spends the saving. (B) is the same GPU and the same weights one stack down: the additive constant moves 4.17 $\rightarrow$ 23 ms, and the 1.7B $\rightarrow$ 0.6B signal falls from 1.32 ms on 0.10 ms of device noise to 0.32 on 1.29. The slope is physics; the intercept is a stack, and it is on no spec sheet.

EVERY OBVIOUS CLOCK TEST FAILS

what the premise seems to promise	what 126 timing cells say
Fewer bytes $\rightarrow$ proportionally less time	4-bit returns 14% of it, and is slower at $B \geq 4$
Batching amortizes the weight read away	the gap is flat (1.32 $\rightarrow$ 1.37 ms); the noise grows 8 $\times$
Half the KV bytes $\rightarrow$ half the slope	a real int4 cache is 1.7 $\times$ steeper
Jitter is positive, so take the minimum gap	73% of an honest server's gaps are 0 ms

Table 3 clock_channel.json, _arch.json and _kvq.json (126 configurations, a real NF4 provider, a real quanto int4 cache, real speculative decoding). The premise survives every row; the naive **absolute** test survives none.

THE TEST THAT SURVIVES: SUBTRACT




Figure 6 slope_verifier_window.json, 3,072 timed decode steps per cell, scored through the **same** harness.evaluate as every token verifier. (A) is measurement only: the subtraction cancels Figure 5's stack constant, and hiding from it costs the 53 ms per token that truncation saved. (B) adds the column's one modelled input, $\sigma = 50$ ms of client-side wire jitter — pessimistic, and one-sided: a busy neighbour only pushes D up, toward the honest side, so co-tenancy costs power, never a false accusation. Read at prefll rather than in the gap the same premise is 25,000 $\times$ cheaper again on the substitution row (clock_algos.json).

WHAT AN AUDITOR ACTUALLY GETS

A price, not a ranking: every verdict costed in tokens $\times$ seconds — and 14 of 35 cells cost infinity.

One protocol, one grid: 8 detectors $\times$ 6 deviations; under an enforced pool ceiling, 16 of 24 cells fall.

A second channel, on the same board: 36–152 tokens of stream, no audit FLOPs, no cooperation.

None of this is a treaty regime — but it prices one. A clause naming a detector cannot be enforced; one naming a pool size, a false-alarm rate and a probe pair can be. The gaps are honest: a fleet over months is not a stream over seconds; the wire is modelled, the device measured; models are 0.13B–8B; no provider here adapts.